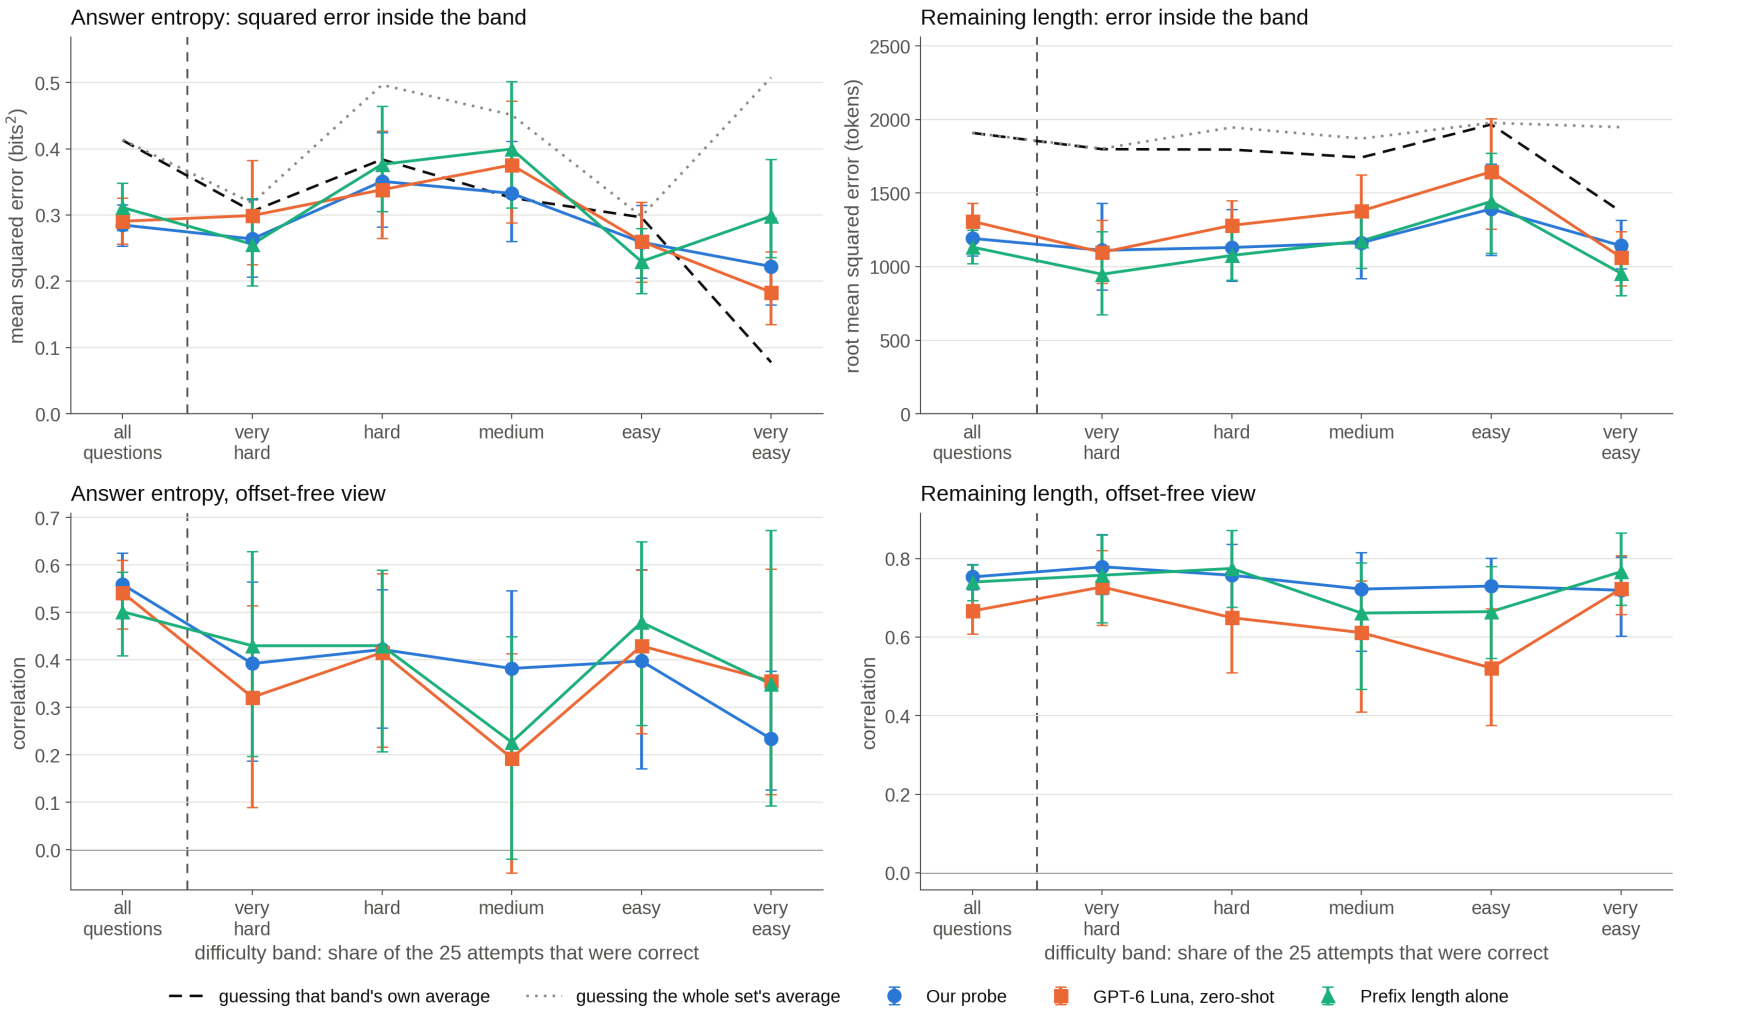
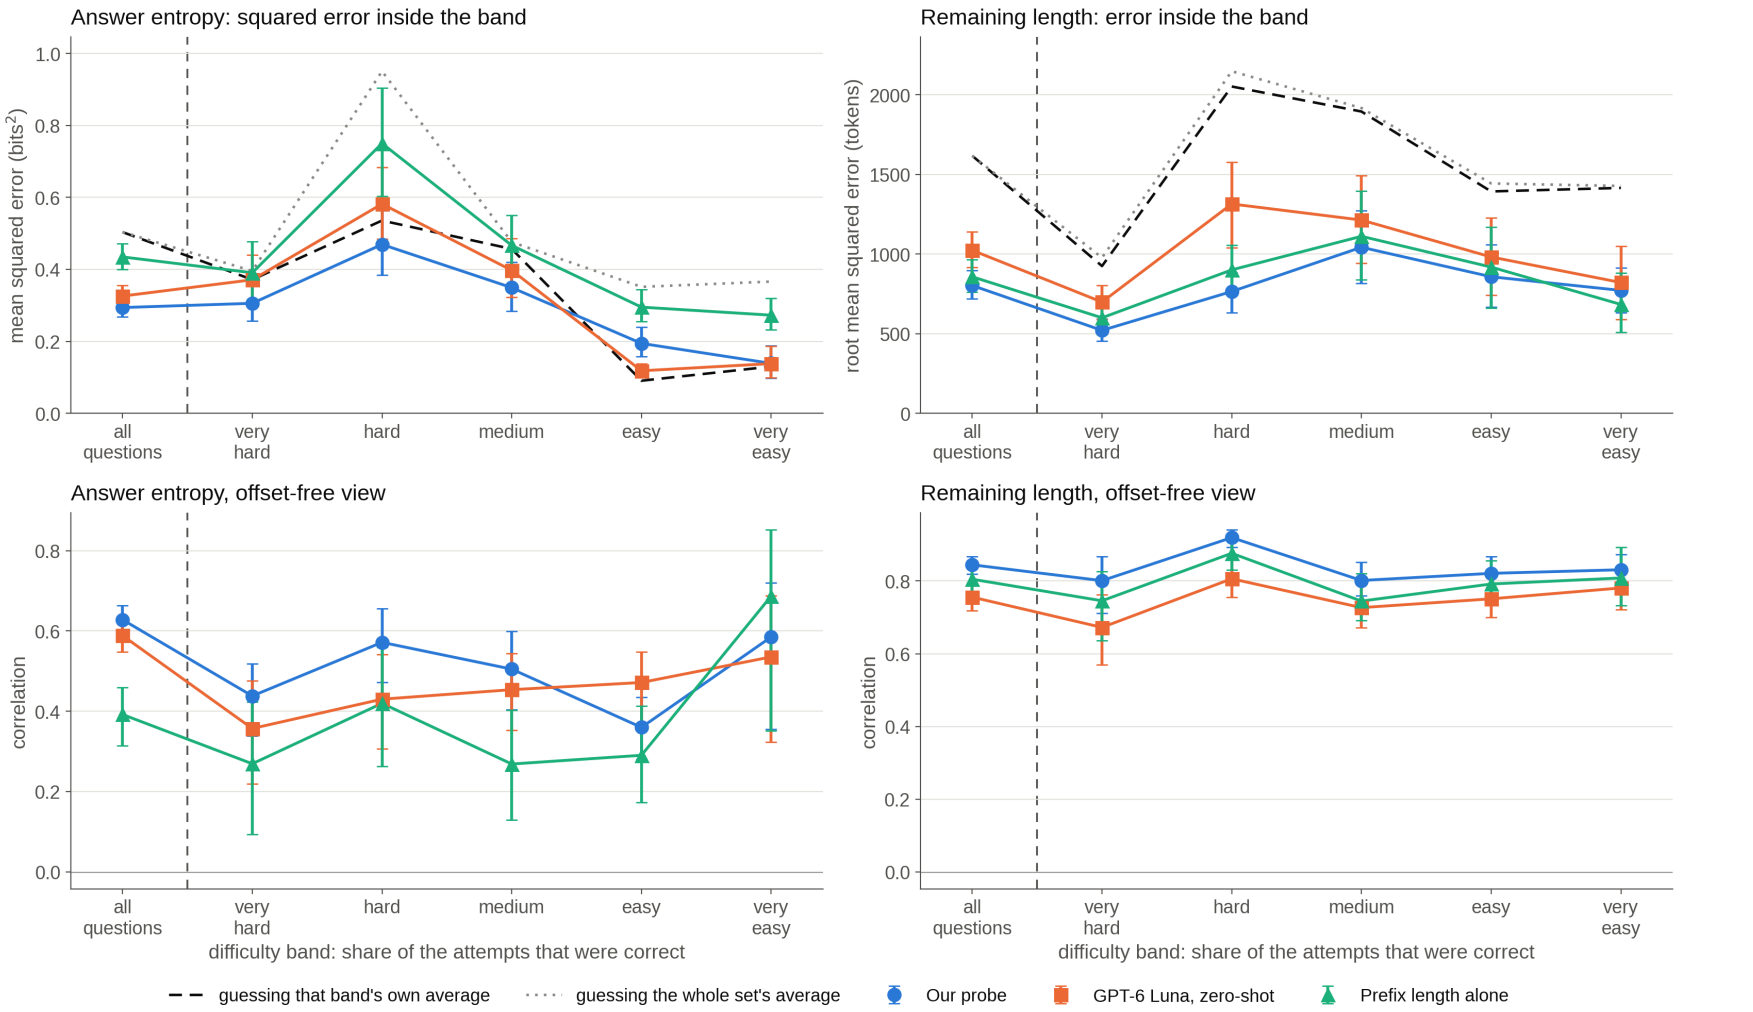
MSE per difficulty band, GPQA. R-squared inside a band divides by that band's own variance, which is the only reason very-easy read -1.090. On MSE it has the lowest error of any band.

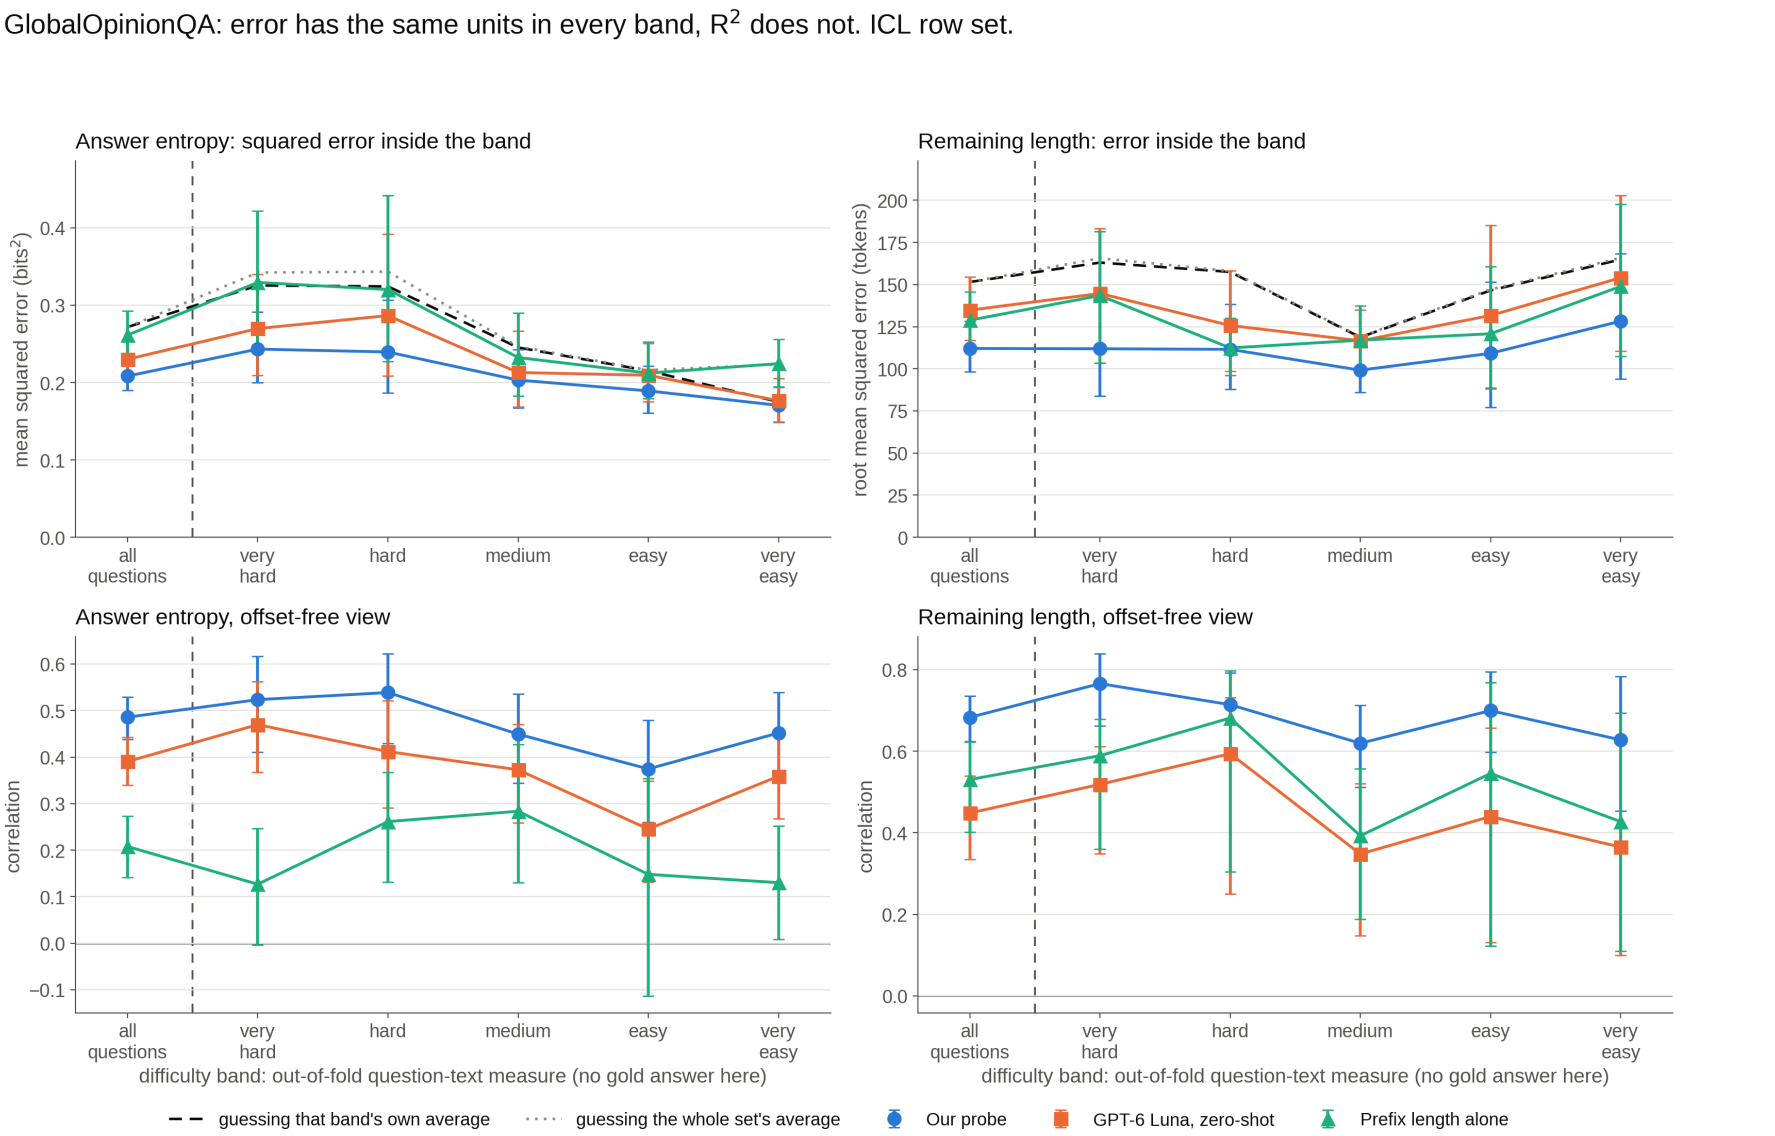


MSE per difficulty band, MMLU-Pro. Same at 12.6x the questions. With 5 rollouts per question, 65 percent sit at accuracy exactly 0 or 1, so read the three middle bands.

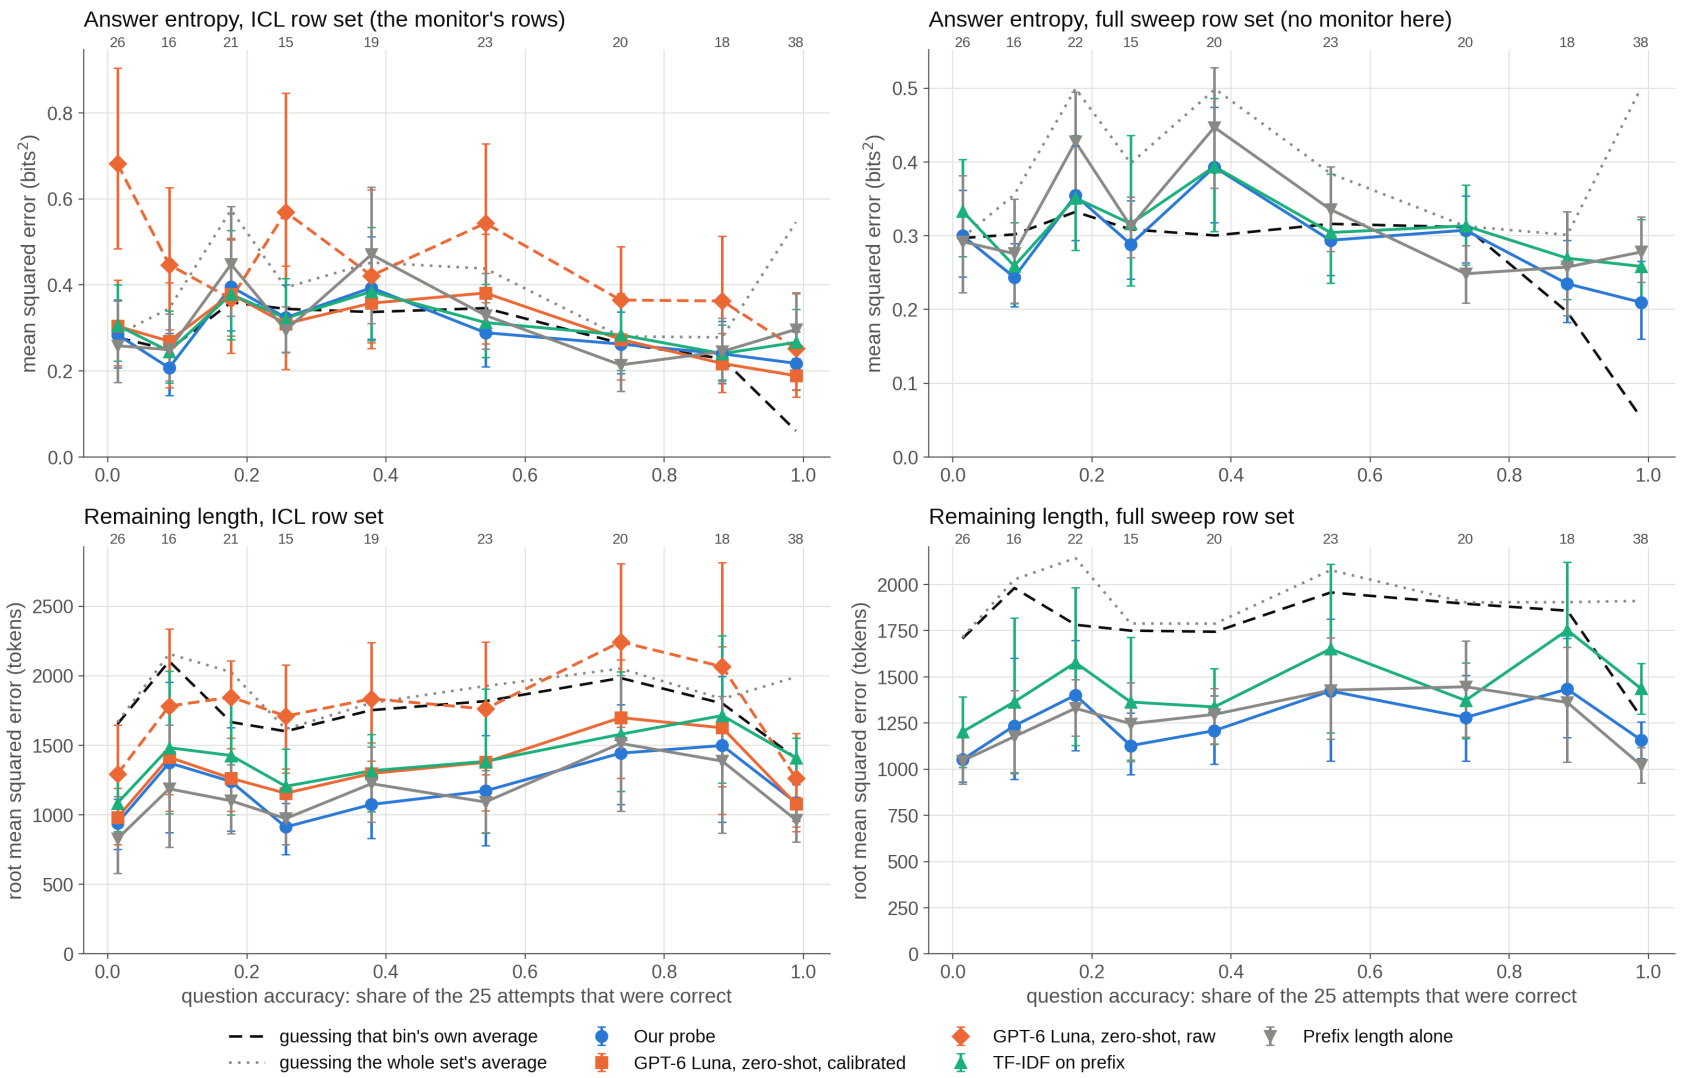
MMLU-Pro: error has the same units in every band, R^2 does not. ICL row set.



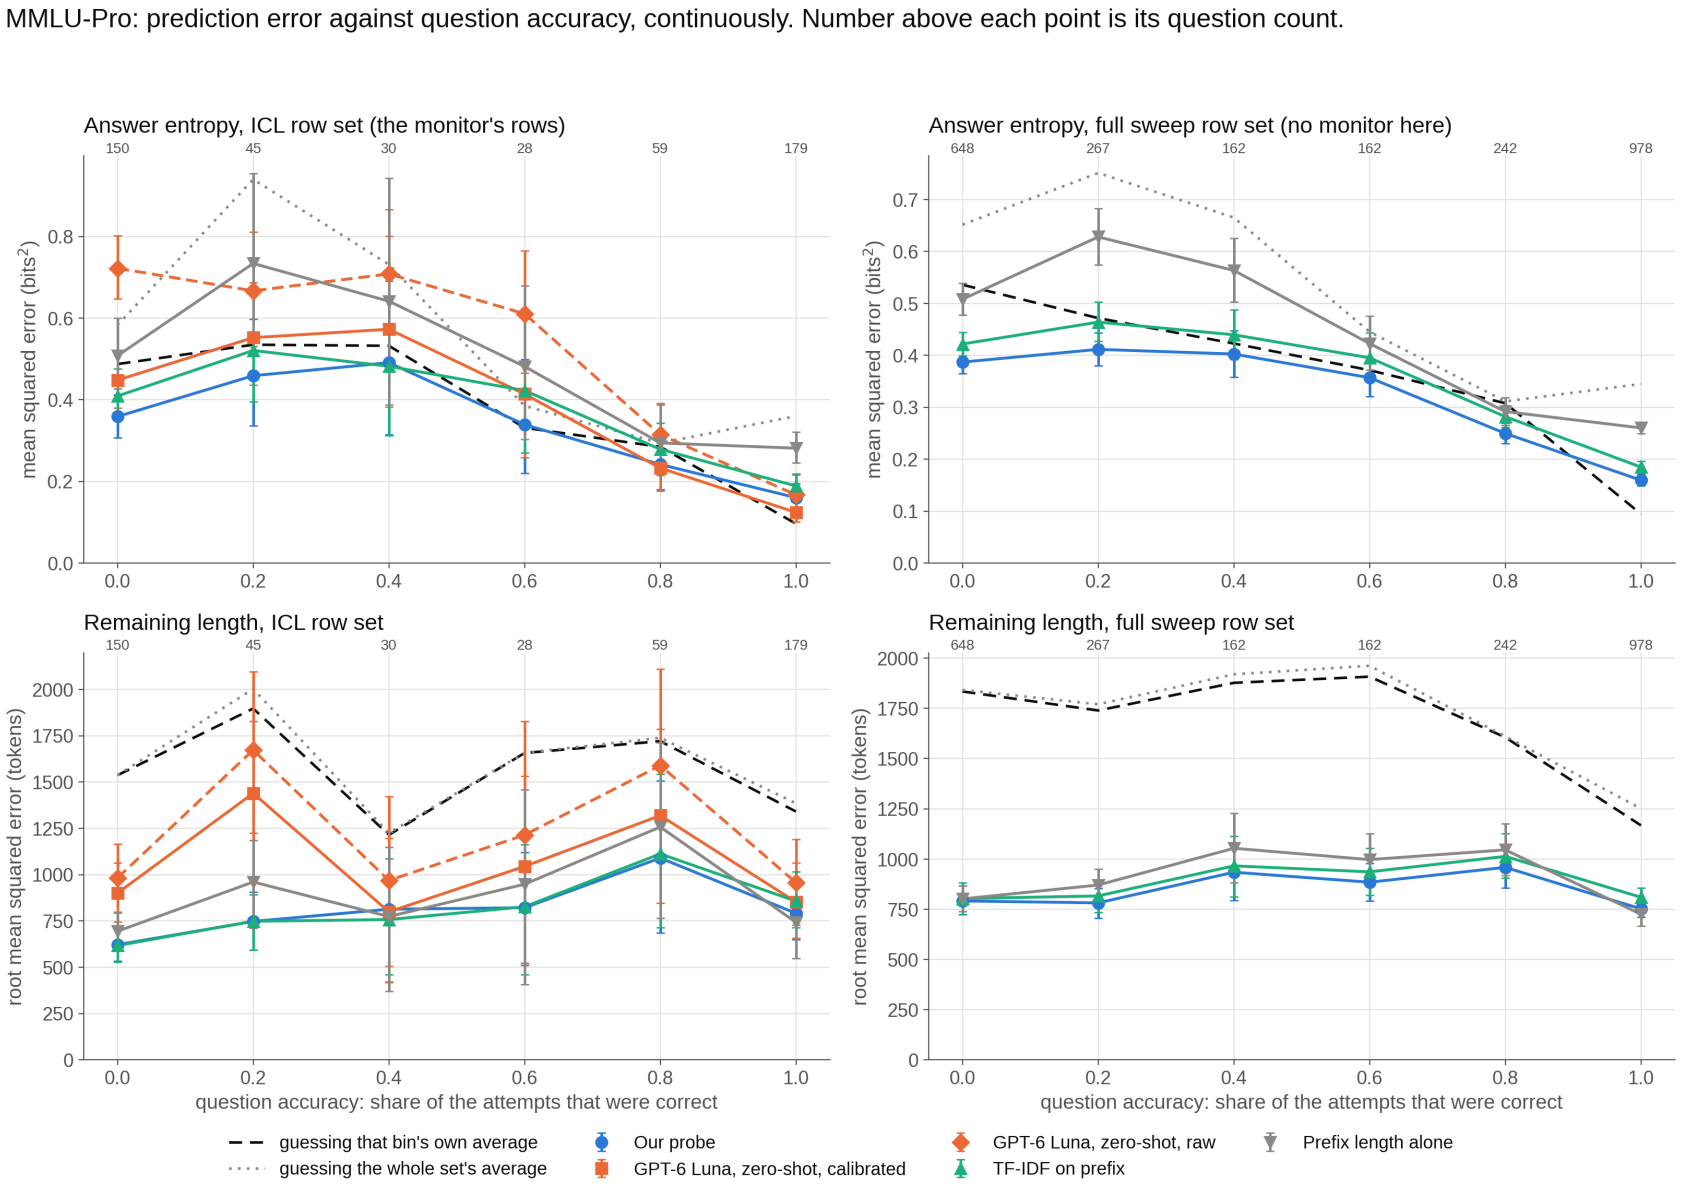
MSE per difficulty band, GlobalOpinionQA. The probe beats the band average in all five bands, and no band is degenerate. No gold answer here, so bands come from a non-circular text-derived measure.



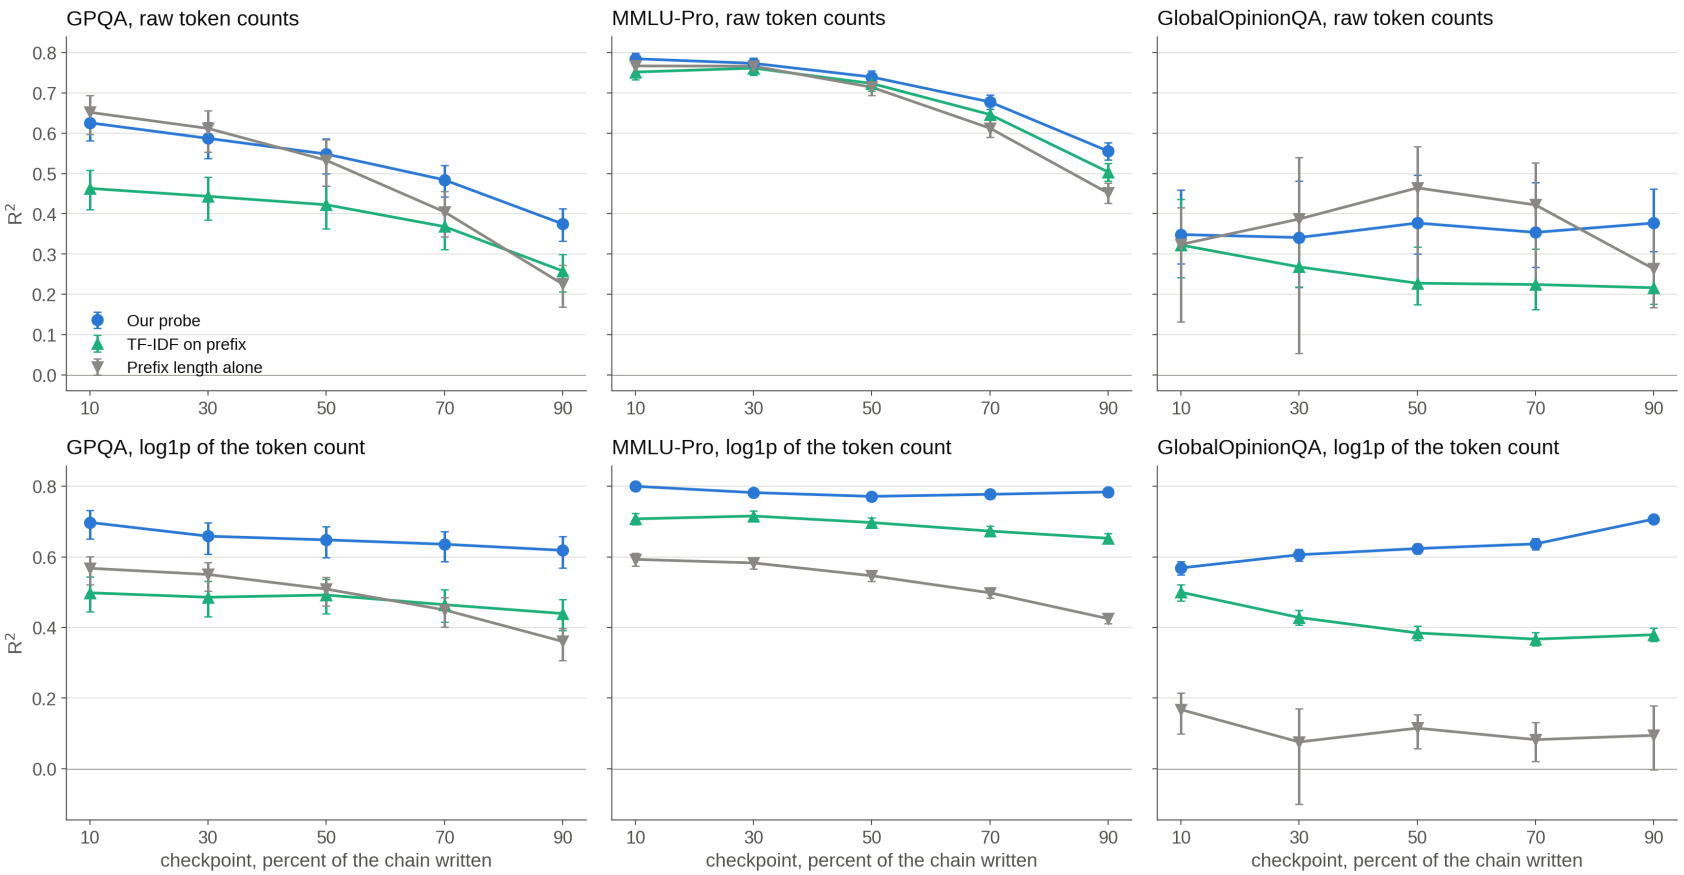
Error against question accuracy, GPQA. Continuous, no buckets. Error falls as accuracy rises. 198 questions, so bins are thin and intervals wide.



Error against question accuracy, MMLU-Pro. Entropy error falls 0.387 to 0.160 as accuracy goes 0 to 1. The probe is under the bin's own average in 5 of 6 bins, losing only where every attempt was correct.



Remaining length against checkpoint. Top row raw, bottom log scale; never compare the two. Prefix length alone matches us on the raw scale early and never on the log scale.



Spread of remaining length against checkpoint. On the log scale this rises with depth on all three datasets, where the raw version decays or is flat.

Spread of remaining length against checkpoint depth. Top row raw, bottom row log scale: never compare one to the other. Full sweep row set.

